Web-based Chess

Patrick Lawrence, Trush Patel, Ishmail Koroma, Kevin Codd, Matt DiStefano

Table of Contents

Summary

The task for this project is to come up with an original idea, design and implement the application, and carry out a full stack deployment with CI/CD services on CloudLab. We decided to build a web-based multiplayer chess application. Our goal is to build a cloud-supported application that supports multiple users and concurrent game sessions with the option to play against an AI at varying levels difficulty. Each user should be availed individual performance statistics to gauge their skill level.
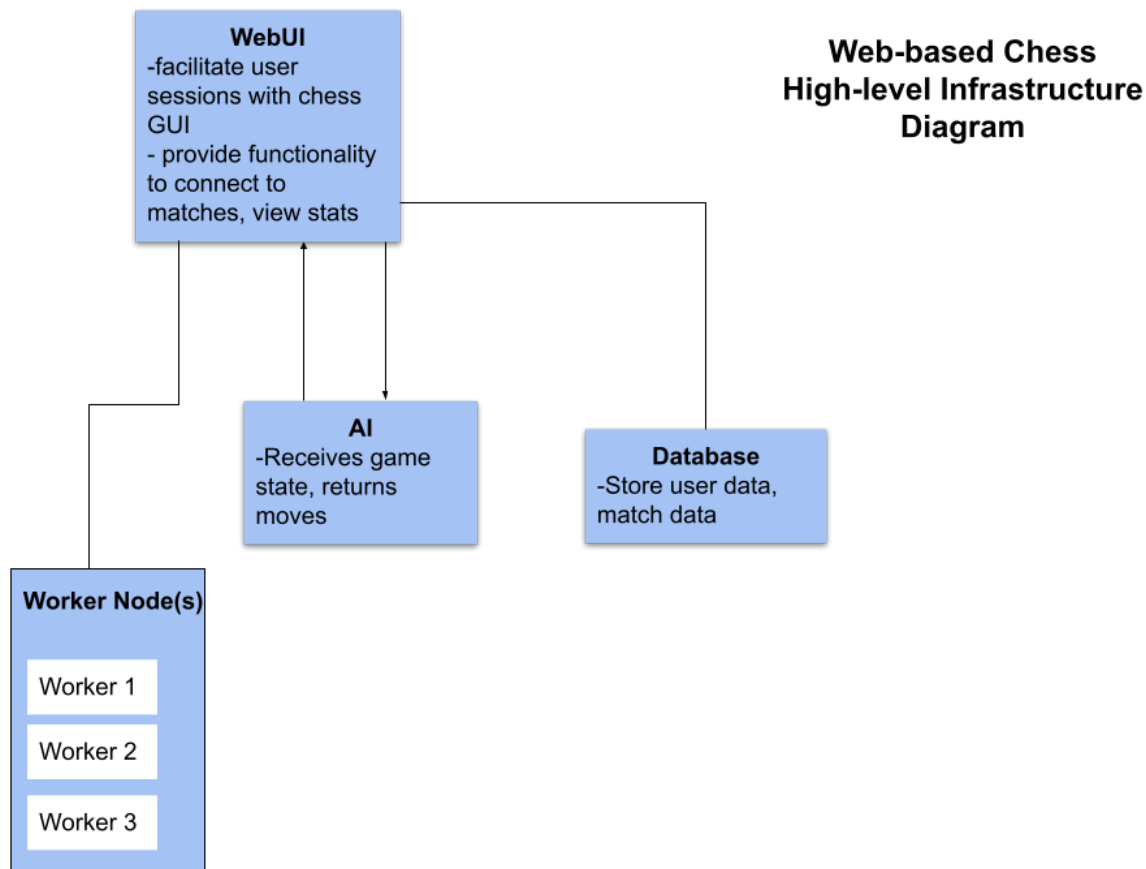
Team Description

We are a team of 5 computer science undergraduate students at West Chester University of Pennsylvania designing, implementing, and deploying a full-stack project for the Introduction to Cloud Computing course. We all come from various backgrounds that include software design as well as front- and back-end development. Each team member's individual qualifications are specified in the Resumes section.

Technical Report

Chapter 1 - Project Vision

Our overall vision for this project is to create a cloud-based chess game playable in the browser — essentially like a simplified version of chess.com or lichess.org. Core functionality would include a fullyimplemented GUI-based chess game with options to play against a human or AI opponent, creation of user accounts and persistent storage of user data, and matchmaking between users. We could use Jenkins for continuous integration and continuous deployment to prevent downtime and automate tasks when making updates to the application. Our primary revenue source will be from ads, i.e. banner ads placed throughout the user interface. We are also considering a wager feature that would bring in revenue by taking a percentage of deposits and/or wagers on games by the participating players. It is important to note that this project's constraints are limited to the extent of the CloudLab hardware available to us. Below is a high-level diagram of the rough infrastructure we envision.

Web-based Chess High-level Infrastructure Diagram

WebUI
-facilitate user sessions with chess GUI
- provide functionality to connect to matches, view stats

AI
-Receives game state, returns moves

Database
-Store user data, match data

Worker Node(s)

Worker 1

Worker 2

Worker 3

We originally planned on including a "master" node, which would have facilitated player matchmaking and served as a broker between the WebUI sessions, players, and AI. Our revised plan tentatively scraps this component of the architecture for similification, opting to use the database as an intermediary between webUI sessions, and to have the AI code housed in the webUI. We would also potentially have a number of "worker" nodes to provide extra compute to the webUI node when necessary, as it will be responsible for facilitating all webUI and AI

sessions. Our hope would be to essentially have the workers function as one node via virtualization for the purposes of the code we deploy to the webUI.

Chapter 2 - Technical Implementation

Web UI

The web UI's key user-facing functionality will be GUI-based chess. The user will be able to play chess on a graphical chess board, most likely with a point-and-click control paradigm. The web UI will also include pages for login and account creation, and for viewing user stats. Additionally, it will serve to connect the user with the revenue framework we choose to implement, such as ads, a wagering system, or a premium subscription service. Some candidate frameworks for implementation include JavaFX, React.js, and Vue.js, but we must consider which will provide the most efficient rendering and convenience for the user.

Key web UI Features

User Actions	Potential Implementation
Create/update/delete account	When a user first connects to a webUI session, there will be a button linking to a sign-up window where the user can create a username and password. We can also have a "user settings" button somewhere in the primary window that brings up options to update username/password or delete an account.
Login	When a user first connects to the web UI, they will be presented with a window with options to enter their username and password or create an account via a button.
View stats	A button will link to a window where a user can pull up basic stats such as win/loss ratio, and their game history. This data would be read and calculated from the database.
Create invitations and join matches	Users can choose to create an invitation - either public or to a specific user. A list of invitations is

	then periodically published to all connected web UI sessions with permission to view them, and other users can accept these invitations. The user can also choose to create a game against an AI rather than play against other users. For all of these options, the user must also specify the starting amount of time for the match timers. (2, 5, 10 minutes)
Move pieces via point-and click/drag and drop	In the gameboard GUI, users can move pieces by dragging and dropping. A background rulechecking function will verify that the move is valid and allow it to be executed on the GUI if it is legal.
Game clocks	Clocks for both players will display how much time they have left. Time controls can be customized, perhaps via a drop-down when game invitations are first created.
Resign/leave a match	There will be a button to resign/leave a match, ending the game session in a win for the resigner's opponent.

Backend Functionality/Database

The database will store user account information including username, rating and game history. It will also serve as a broker to facilitate game invitations and the creation of game sessions. The database could be implemented using several different database frameworks, including mySQL, Redis, and MongoDB. However, we are leaning toward PostgreSQL because it is an object-relational database which makes it easier to instantiate and store game sessions. This could be useful for searching game history by player, interacting with the game AI, and restoring crashed sessions. Further, PostgreSQL is better-performing for concurrent reading/writing than our main alternatives (2), which could be critical for scaling to facilitate many matches at once.

To create matches, users will be able to send public invites to other users, or private invites to specific users. In our current architectural plan, these invites will be written to the database, perhaps with critical information about how to connect a game, such as the port number of the web UI session originating the invite. Web UI sessions will periodically query a list of invites from the database, which will refresh in the user's GUI. This system would be able to facilitate matches between several web UI, with the main constraint being the compute resources allocated to running web UI sessions (and AI players).

To facilitate games, we could potentially take a few different approaches. First, one web UI session could

"host" the game, storing the moves and game timer in memory, receiving moves from the opponent's web UI session, and also pushing updates to the opponent. We could also potentially have the database server as the host, with every move and the game clock being written to it while the game is in progress. The viability of this approach would largely depend on latency. Finally, if we find that these approaches are infeasible or difficult to implement, we could switch to an event-driven networking model, utilizing a "master" node, as in our original architecture. The master would facilitate game sessions between users and store game state independently of web UI sessions, like the database as described above, but in the node's memory. The master would also track connected web UI users and game invites in memory, and be an alternative to serve as the "host".

When a game is finished, the game will be saved to persistent storage in the database, and player rating changes will be computed, likely in the WebUI backend. We currently also intend to house the AI code on the webUI backend.

To provide the compute resources necessary to run multiple webUI sessions and AI players, we intend to allocate a number of "worker" nodes to the webUI. Ideally, we could combine multiple hardware nodes into one virtual machine for simplicity, perhaps using Kubernetes, Linux Virtual Server, or another option easily deployable on Cloudlab.

Most of our backend code could likely be implemented in Java and SQL.

Key Backend Features

Functionality	Potential Implementation
Check that moves are legal before allowing execution.	This legality check could run on the WebUI backend, which would be optimal for reducing lag. A function would take in the current game state and proposed move and ensure it is in compliance with all rules.
Send moves to database, opponent webUI session, AI.	Once a legal move is executed, it needs to be sent to the opponent through one of the means described above.
Send and receive game invites.	When a user creates a game invite, it must be pushed to all other users with viewing permissions. Similarly, the WebUI must be open to read/receive a list of open invites. Our current plan is to read/write these lists from the database. Entries could include fields specifying viewing permissions and webUI port numbers.

Keep registry of users connected to web UI sessions	The database will keep track of all users that have connected to the network with their login as well as the players that are connected and currently playing a game.
Push list of game invitations to users	When a game invitation is created, it will be
	written to a registry in the database, which other webUI sessions will regularly query. If the invite is public, all users will see the invitation. If the invite is to a specific user, only that user will see the invitation.
Track game state for in-progress games	The sequence of moves made so far in game will be written to the database in real-time, stored in the memory of a webUI session, or stored in the memory of a "master" node as in an event-driven model.
Timeout/auto-resign	Should a user have connection issues or disconnect from the application, there will be a small amount of time to allow the user to reconnect and resume play. When logging back in, the database will have a record of if the user is supposed to be in a game. The webUI can query this record upon login, and automatically connect the user if they are supposed to be in an inprogress game. Should the user not connect in time, they will automatically resign.
Initialize a game	When a game between two players is created via an accepted invite or a request to play against the AI, colors will be randomly assigned and the clock will be started.
Timers	The official time will be stored in the database or host WebUI session and pushed to players on a regular basis.
Write results and updated player stats to database	At the end of the game, the game could be encoded in a file format such as FEN or PGN and written to the database. Players' game history could also be updated.

Update player ratings (tentative)	We could implement an elo rating system where at the end of every game the master calculates rating updates and writes them to the database.

AI Chess Player

We'd like the user to have the option to play against an AI chess player should they want to. Implementing an AI is an important stepping stone for us in this project because it's very likely that there will not be any users to matchmake with at times. Having the AI player will allow these users to use the site during these slower times.

Behind the scenes, this AI will receive a game state and return what it believes to be the best possible move for each board position. A feature like this could be implemented in a couple different ways. We could bite the bullet and implement our own AI using the well-known minimax algorithm. This is a recursive or backtracking algorithm used in decision-making and game theory. It provides an optimal move for the player assuming that the opponent is also playing optimally. Alternatively, our other option is utilize open source options that already exist such as Stockfish to do this task for us.

Stockfish is one of the most well-known chess engines available to the public and it's considered one of the strongest as well. Many chess players consider the engine to be stronger than human players at the game. It would be able to provide better moves for board positions than our algorithms could and it would greatly increase the quality of the user experience.

In either case we would need to implement fine tuning the AI difficulty to allow users of all skill levels to play against the AI. This would likely involve giving it a probability of playing varying degrees of suboptimal moves.

We are currently leaning towards using Stockfish, as creating our own AI would be way too timeintensive for this project. We've decided to consider adding our own AI algorithms as a "stretch" goal for us to experiment with once other key functionality is complete.

In terms of deployment, we currently plan on hosting the AI in the webUI backend.

Chapter 3 – Intermediate Milestones

As of April 17, we have implemented the foundation of our web application's infrastructure. The front-end consists of a chess user interface found on GitHub. Since we have also built out all the supporting files for the Kubernetes clusters, we are able to get the game interface up, running, and functional on the Kubernetes branch of our project. Additionally, the Stockfish AI image and Dockerfile are built out.

We have run into a number of challenges with getting the front-end to work and in building out the continuous integration and continuous delivery with Jenkins. The user interface only worked on the Kubernetes branch if each command was executed one-by-one. When we ran a bash script, the same error

consistently occurred when pushing the docker images: connection to port 30000 refused. We addressed this by adding sleep commands in the shell script to allow time for longer commands to execute. Challenges with the CI/CD revolve around configuring Jenkins to support our project's specifications. We found that the Golang server template provided requires significant modifications to accommodate our multi-pod Kubernetes cluster.

Our primary goal looking forward is to get the CI/CD pipeline functional with our project's specifications, allowing us to more rapidly and efficiently iterate in our development cycle. Next in priority is building out front-end features such as a menu system, and supporting chess games between multiple users and/or an AI agent. Once we have the barebones functionality of the application established, we can move onto more intricate features such as game clocks, pages for viewing players stats, ad placement, and other bells and whistles.

Kubernetes Cluster

As of this deliverable, we've been successful in running our application on a Kubernetes cluster deployed on Cloudlab, and accessing it through our head node URL. We were able to implement Kuberntes .yml files for deployment of our pods and services by following the format in ram_coin. We were able to do the same for our docker-compose files. Although our app has three pods, docker-compose and Kubernetes allow for build and deployment with empty docker files, which we have leveraged to test the webUI without our other pods being ready for deployment yet. As the AI is not quite ready for integration with the front-end, we have not yet included it in our configuration files. Further, we have not yet implemented the database, so its Dockerfile is currently blank. The webUI is fully configured for the Kubernetes cluster. We were able to borrow and make slight modifications to the ram_coin Docker and Kubernetes configuration files for the webUI, as ours is also a Node.js application. We were first able to successfully test our cluster by following the commands in the Kubernetes lecture, replacing ram_coin's git repo and Kubernetes services with ours, and logging into the webui port on the Cloudlab head node. After successfully viewing a message we printed to the console, we pointed the web UI at an index.html file, and were able to begin building it out from there. We have since converted the commands necessary to deploy the cluster to a bash script.

Jenkins CI/CD

Our work on configuring Jenkins for our application has mostly consisted of adapting the Jenkinsfile for the example golang server from "Kubernetes Application: CI/CD pipeline - Part II" to suit our Kubernetes cluster, as well as the pursuant changes to the setup in the Jenkins GUI. We were successful in deploying the application with Jenkins by porting over the commands from the aforementioned bash script for launching the cluster into the Jenkinsfile. However, this deployment only works once – once the services are up and running we don't yet do anything to reset them, thus running another Jenkins build tries to over-write our already up-and-running services, causing an error. We will have to adapt our commands and overall pipeline configuration to update the services with new source code rather than attempting to re-initialize them without properly killing them.

WebUI

The user interface currently consists of a simple chess board with moving pieces, valid moves, and move suggestions. It is a node.JS application that we found here. Since an existing framework exists that has the essential functionality we require for a chess game, there was no need for us to build one from scratch. Once the database and AI are built, we will be able to post and get information to and from each container. Our goal is to add more features once the infrastructure is completely built out (i.e., game clocks, player stats, etc.). We would like the user to start at a login page, jump to a dashboard where they can create new game sessions and observe their performance statistics, and finally access a page for the live game session.

AI

We currently have a Stockfish Docker image set up and a dedicated container for it. We've automated the compilation of the Stockfish source code and made it seamlessly run Stockfish once compiled. Although, right now we're working on a solution that will allow us to pipe input into the Stockfish console and a solution that will allow us to pipe the output to another one of our containers on our Docker network. We believe the most efficient way to fix this issue is to use HTTP get / post commands to communicate between containers. Once we've resolved this issue we can continue to add features that we originally wanted to implement. Features such as AI difficulty tuning and saving completed games to a database. Overall, we're making swift progress with Stockfish and are continuing to keep it a part of our project.

Database

We have yet to configure the PostgreSQL database. Our current webUI framework stores game states in memory, so the database will tentatively only be necessary for persistent storage of games and player stats. Once implemented, we plan on storing players' game history, credentials, and elo rating. The Web UI will query and write to the database using HTTP get/post.

References

1. https://github.com/trushpatel/CSC468Group5.git
2. https://github.com/tudorfis/chess
3. SQLite vs MySQL vs PostgreSQL: A Comparison Of Relational Database Management Syst

Kevin Codd

Data Scientist // Software Engineer

✉ kncodd@gmail.com
📞 +1 (484) 947- 9711
🔘 github.com/KevinCodd
🔗 linkedin.com/in/kevin-codd
📍 Greater Philadelphia

EXPERIENCE

Data Science Intern   Jun 2021 – Present
CellCarta   Remote (Fremont, CA)

- Researched and improved upon multivariate time series outlier detection algorithm for Flow Cytometry data
- Developed data pipelines in Python and R for comparative evaluation of outlier detection algorithms on many datasets
- Contributed to R API toolkit for CellEngine data analysis software

Data Science Intern   Feb 2020 – May 2021
nth Solutions   Exton, PA

- Developed Java backend for IMU sensor data visualization application
- Modelled time series data using several Python and MATLAB libraries for state space modelling, deep learning, and frequency domain analysis
- Researched and implemented IMU sensor fusion algorithm for accurate orientation tracking
- Led team of interns responsible for data engineering and analysis tasks

Undergraduate Research Assistant   Spring, Fall 2021
West Chester University   West Chester, PA

- Analyzed break-in attempts to Computer Science department's servers using Linux authentication logs and PySpark in support of academic research paper
- Cleaned, parsed, and analyzed logs using PySpark to identify and quantify attacks
- Currently developing backend for data-streaming dashboard using Spark and Kafka

Regulatory Affairs Intern   May 2018 – Nov 2018
Lungpacer Medical   Exton, PA

- Organized clinical trial data and produced reports analyzing statistical trends
- Conducted scientific literature searches to substantiate efficacy of medical device treatment approach
- Proofread and edited technical documentation to be submitted to regulators

EDUCATION

B.S. in Computer Science & Minor in Mathematics   Expected Dec 2022
West Chester University

- **GPA**: 3.95 / 4.0
- **Relevant Courses:** Big Data Engineering, Mathematical Statistics, Artificial Intelligence, Software Engineering, Data Structures & Algorithms, Computer Systems, Linear Algebra, Multivariable Calculus

SKILLS

- **Languages:** Python, Java, R, C, MATLAB, SQL, SAS
- **Tools/Frameworks:** *Data analysis/engineering*: NumPy, SciPy, Pandas, Matplotlib, PySpark, Kafka. *Machine learning*: Scikit-learn, TensorFlow. *Development*: Git, Jupyter Notebook, Anaconda.

ACTIVITIES

Competitive Programming Club   Aug 2020 – Present
West Chester University

- Club President (Aug 2021-Present) - Facilitated club meetings and competitions, created and curated practice material
- Algorithmic programming contests: member of team advancing to 2021 ICPC North American Division Championships, 3rd place team out of 24 in CCSC Eastern 2020 Programming Competition

Matthew C DiStefano

13mdistefano@gmail.com

805 Windridge Lane | Downingtown | PA | 19335 | tel: 610.247.8772

Education

West Chester University of Pennsylvania

Bachelor of Science in Computer Science

3.64 GPA | Expected to graduate May 2022

Relevant Courses

Computer Systems, Data Structures and Algorithms, Calculus, Database Management Systems, Data Communication and Networking, Edge Computing and Deep Learning, Software Testing, Intro to Cloud Computing, Digital Image Processing

Professional Experience

Law Clerk | June 2014 – August 2017

Dalton & Associates, LLC | Wilmington, DE

– Managed and maintained the systems for closing of files and server storage

– Maintained and kept track of client's medical bills and records in accordance with HIPPA standards for ease of access by attorneys and clients alike

– In charge of maintaining files within our cloud system

Law Clerk | June 2019 - Present

Potamkin ARM, LLC | Downingtown, PA

– Keep over 170 entities up to date and manage documents that are to be kept in the minute books

– Maintained timely communication with fellow employees to better identify and resolve issues

Computer Science Projects

CSC496 Final Project | West Chester University | December 2021

– Lead a team to the creation of a graph based shortest path algorithm used to determine the earliest semester a specified Computer Science course could be taken at WCU

Skills

– Languages: Java, JavaScript, C#, Python, SQL, HTML, OCaml

– Tools/Frameworks: Git, Node.js, React.js

Patrick Lawrence

West Chester, PA
(610)-675-8651
patrickjameslawrence@icloud.com

Education

West Chester University of Pennsylvania
Bachelor of Science in Computer Science | May 2022

- 3.38 Cumulative GPA

Related Experience

IT Helpdesk Consultant | December 2020 – Present
West Chester University of Pennsylvania | West Chester, PA

- Identify user issues and dispatch them to appropriate teams when needed
- Educate fellow consultants on internal IS&T functions and solutions to issues users face
- Maintain the IS&T knowledge base; adding articles to pass down my expertise when needed
- Maintain cordial communications and de-escalate situations with bad-tempered users

Software Engineer Intern | June 2021 – July 2021
Universal Health Services (UHS) | King of Prussia, PA

- Built new FTP connections to transfer patient records, test results, and other data between UHS network hospitals
- Utilized SQL and C# to create applications in .NET
- Maintained adherence to all HIPPA privacy standards and protected patient data
- Shadowed UHS software engineers in their meetings and projects
- Reviewed standard UHS operating procedures, all internal IS&T teams and their functions, and the software development lifecycle
- Learned office etiquette and best practices from other employees in the office

Skills Summary

- *Programming Languages:* Java, C++, JavaScript, C#, SQL
- *Development Tools:* Version control & bug tracking
- Heavy interest in advanced mathematics

Projects

Synth | September 2020 – March 2021
Multi-function Discord bot | Node.js and MySQL

- Created a robust command and event handler
- Designed a local MySQL database to store various configuration information from servers and users

Keychain | October 2020
Social media sharing service | Android Studio and MongoDB

- Designed an Android application using Android Studio
- Utilized database services to create a secure account login solution

Trush Patel

(267)-261-9675 | trushp3@gmail.com | linkedin.com/in/trushpatel | github.com/trushpatel

EDUCATION

Georgia Institute of Technology   Atlanta, GA
Master of Science in Computer Science   *Expected August 2024*

West Chester University of Pennsylvania   West Chester, PA
Bachelor of Science in Computer Science   *May 2022*
  - GPA: 3.99/4.00

EXPERIENCE

Software Engineer Intern - Product Team   February 2022 – Present
humanID   *New York, NY*
  - Developing anonymous authentication solutions for secure single sign-on
  - Leading Scrum teams to consistently deliver new features and upgrades
  - Constructed concise documentation to accelerate client integration by 50%
  - Building a developer dashboard to manage client integrations and aggregate relevant data

Tutor - Learning Assistance & Resource Center   December 2021 – May 2022
West Chester University of Pennsylvania   *West Chester, PA*
  - Organized one-on-one meetings to assess comprehension and teach fundamental concepts
  - Instructed the development of effective time management and test preparation strategies
  - Tutored undergraduate students in Introduction to Statistics, Foundations of Computer Science, and Data Structures and Algorithms
  - Acquired CRLA Level I certification for professional tutors and peer educators

Undergraduate Research Assistant - Dr. Jongwook Kim   October 2021 – May 2022
West Chester University of Pennsylvania   *West Chester, PA*
  - Researched software design patterns and refactoring paradigms
  - Assessed program transformations that augment software reliability and extensibility
  - Conducted tests to determine the efficacy of refactoring tools in programs with 10 to 10,000 lines of code
  - Refactored a program that reduces the execution time of existing refactoring tools by 35%

PROJECTS

KubeChess | *PostgreSQL, Express.js, React.js, Node.js, Docker, Kubernetes, Jenkins, Ansible*   May 2022
  - Collaborated with 4 developers to build an online playground for chess players
  - Implemented a React.js front-end and RESTful backend that uses Express.js, Node.js, and PostgreSQL
  - Deploying with Kubernetes to scale for multiple users and concurrent sessions
  - Implementing a Jenkins pipeline to support CI/CD

Yelp Dataset | *Python, SQL, Apache Spark*   December 2021
  - Built an algorithm that implements PageRank to identify influencers in a 10-gigabyte dataset
  - Computes the top 5% of influential users from quantity and quality of reviews and connections
  - Extracted each user's review data to identify individual preference in type of places reviewed
  - Pinpointed reviewed business locations to cluster users geographically

HONORS & AWARDS

Upsilon Pi Epsilon | *International Honor Society for the Computing and Information Disciplines*   May 2022
3rd Place | *PACISE Regional Programming Competition*   December 2021
Dean's List | *West Chester University of Pennsylvania*   December 2020 - December 2021

TECHNICAL SKILLS

Languages: Java, Python, C/C++, JavaScript, HTML, CSS, OCaml, Bash, SQL
Tools/Frameworks: Apache Spark, Bootstrap, Django, Docker, Git, Jenkins, Kubernetes, Node.js, React.js
Coursework: Artificial Intelligence, Big Data Engineering, Computer Systems, Data Structures and Algorithms, Introduction to Cloud Computing, Linear Algebra, User Interfaces, Software Engineering, Software Testing

Morlai Ishmail Koroma

106 Norma Road, Yeadon PA 19050   ishmailkoroma1@gmail.com (267) 515 2702

EDUCATION

West Chester University   Expected Graduation Date: May 2022
700 S High St, West Chester, PA 19383
Major: Computer Science

TECHNICAL SKILLS:

- Proficient in JavaScript, C++, Python, HTML, Linux, and Microsoft Office

PROFESSIONAL EXPERIENCE:

U.S. Navy   **Naval Research Enterprise Intern**
Philadelphia, PA   October 2021 – December 2021
- Researched and provided recommendations and new methods for increased monitoring, reducing maintenance, and repair of critical naval machinery systems on Navy ships.
- Provided a possible application of LoRa sensors on the Lube Oil System on Navy ships to detect water leaks, maintain temperature pressure and reduce the maintenance cycle for sailors.
- Developed a cost analysis framework for the materials used in the implementation of the LoRa sensors on the Lube Oil System.

U.S. Navy   **Naval Research Enterprise Intern**
Philadelphia, PA   June 2020 – August 2020
- Researched methods to enforce LoRa (long-range) technologies on Naval platforms without compromising security.
- Assessed the LoRaWAN (long-range, wide area network) vulnerabilities and cyber security risks involved in implementing long range sensors on Naval platforms

Avista Healthcare   **Dietary Aide**
Cherry Hill, NJ   August 2019 – March 2020
- Monitored inventory and stock for kitchen ingredients
- Maintained kitchen equipment and appliances via guidelines of the facility's dietary manager and kitchen cook
- Oversaw take-down of dining areas, collecting all silverware, discarding of leftovers and removal of garbage and recycling items.

Sky is the Limit (Moving Company)   **Moving Assistant**
Greater Philadelphia   June 2018 – May 2019
- Tracked payments for services
- Kept detailed inventory of customer items to ensure safekeeping and integrity of items moved.
- Kept track of communication and interaction with all customers via emails, texts, and calls

COMPUTER SCIENCE PROJECTS:

CSC 472 Final Lab Project   **Lead Member**
West Chester University   November 2021 – December 2021
- Performed a successful multi-stage exploit attack on a file to get and reveal content inside the target flag.
- Completed format string, ROP, GOT overwrite, and stack overflow attacks on the program to leak specific information.

CLUBS & VOLUNTEER EXPERIENCE

WCU Weekly   **Videographer/Editor**
West Chester University   August 2020- December 2020
- Filmed highlights at West Chester University sports games and helped edit weekly sport segments

Christ The King Prayer Chapel.   May 2018 to August 2018
Philadelphia, PA
- Filmed services.
- Helped organized the church and collected donations/offerings from members

INTERESTS:

- Cybersecurity; Network and Computer Systems; Artificial Intelligence; Computer Graphics; Video Games; Basketball; Football